

oneAPI programming in Julia with oneAPI.jl

Tim Besard

Why Julia?

High-level programming language
designed for performance

Why Julia?

High-level programming language
designed for performance

```
julia> data = (1, rand())  
(1, 0.5326182923218289)
```

```
julia> sum(data)  
1.532618292321829
```

Why Julia?

High-level programming language designed for performance

```
julia> data = (1, rand())  
(1, 0.5326182923218289)
```

```
julia> sum(data)  
1.532618292321829
```

```
julia> @code_llvm sum(data)  
define double @julia_sum({ i64, double }* nocapture nonnull readonly align 8 dereferenceable(16) %0) #0 {  
top:  
  %1 = getelementptr inbounds { i64, double }, { i64, double }* %0, i64 0, i32 0  
  %2 = getelementptr inbounds { i64, double }, { i64, double }* %0, i64 0, i32 1  
  %3 = load i64, i64* %1, align 8  
  %4 = sitofp i64 %3 to double  
  %5 = load double, double* %2, align 8  
  %6 = fadd double %5, %4  
  ret double %6  
}
```

Why Julia?

High-level programming language designed for performance

```
julia> data = (1, rand())  
(1, 0.5326182923218289)
```


```
julia> sum(data)  
1.532618292321829
```

```
julia> @code_native debuginfo=:none sum(data)  
_julia_sum:                               ; @julia_sum  
; %bb.0:                                   ; %top  
    ldp    d0, d1, [x0]  
    scvtf  d0, d0  
    fadd   d0, d1, d0  
    ldr    d1, [x0, #16]  
    fadd   d0, d0, d1  
    ret
```

GPU support in Julia

GPU-enabled applications

- Flux.jl (deep learning)
- CLiMA (ocean modeling)
- DifferentialEquations.jl
- Yao.jl (quantum information)
- ...



oneAPI.jl



Metal.jl



AMDGPU.jl



CUDA.jl

Shared infrastructure

- GPUCompiler.jl
- GPUArrays.jl
- KernelAbstractions.jl
- ...

Easy to get started

1. Download and unpack Julia 1.8: <https://julialang.org/downloads/>

2. Launch Julia and enter the the package manager

```
pkg> add oneAPI
```

3. Import and verify the oneAPI.jl package

```
julia> using oneAPI
```

```
Downloading artifacts: ...
```

```
julia> oneAPI.versioninfo()
```

```
Binary dependencies:
```

```
- NEO_jll: 22.53.25242+0
```

```
- libigc_jll: 1.0.12812+0
```

```
- ...
```

```
1 device:
```

```
- Intel(R) Arc(TM) A770 Graphics [0x56a0]
```

Automatic download of
binary dependencies

Array abstraction

```
julia> vec = oneArray([1])  
1-element oneVector{Int64, oneL0.DeviceBuffer}:  
 1
```

`oneArray` serves multiple purposes:

1. container for device memory
2. abstraction for data-parallel programming

```
julia> vec .+ 1  
1-element oneVector{Int64, oneL0.DeviceBuffer}:  
 2
```


Array abstraction

Linear algebra:

```
julia> using LinearAlgebra
```

```
julia> vec = oneVector(rand(Float32, 2))  
dot(vec, vec)
```

```
julia> mat = oneMatrix(rand(Float32, 2, 2))  
mat * mat
```

Higher-order functions:

```
julia> map(vec) do val  
    val + 1  
end
```

```
julia> reduce(+, vec)
```

Obviates kernel
programming!

Statistics:

```
julia> using Statistics
```

```
julia> mean(mat)
```

```
julia> std(mat)
```

Kernel programming

```
function vadd(a, b, c)
    function kernel(d_a, d_b, d_c)
        i = get_global_id()
        d_c[i] = d_a[i] + d_b[i]
    return
    end

    d_a = oneArray(a)
    d_b = oneArray(b)
    d_c = oneArray(c)

    len = prod(size(a))
    @oneapi items=len kernel(d_a, d_b, d_c)
    c .= Array(d_c)
end
```

Similar to CUDA.jl, AMDGPU.jl, ...

Differences with DPC++/SYCL

- OpenCL intrinsics
- Global semantics

Kernel programming

```
function vadd(a, b, c)
    function kernel(d_a, d_b, d_c)
        i = get_global_id()
        d_c[i] = d_a[i] + d_b[i]
        return
    end

    d_a = oneArray(a)
    d_b = oneArray(b)
    d_c = oneArray(c)

    len = prod(size(a))
    @oneapi items=len kernel(d_a, d_b, d_c)
    c .= Array(d_c)
end
```

```
julia> @device_code_llvm vadd([1], [2], [0])

define spir_kernel void @kernel(
    { { [1 x i64], i8 addrspace(1)* } }* byval,
    { { [1 x i64], i8 addrspace(1)* } }* byval,
    { { [1 x i64], i8 addrspace(1)* } }* byval
) local_unnamed_addr {
entry:
    %3 = call i64 @_Z13get_global_idj(i32 0)
    ...
    store i64 %14, i64 addrspace(1)* %18, align 8
    ret void
}
```

Kernel programming

```
function vadd(a, b, c)
    function kernel(d_a, d_b, d_c)
        i = get_global_id()
        d_c[i] = d_a[i] + d_b[i]
    return
end

d_a = oneArray(a)
d_b = oneArray(b)
d_c = oneArray(c)

len = prod(size(a))
@oneapi items=len kernel(d_a, d_b, d_c)
c .= Array(d_c)
end
```

```
julia> @device_code_spirv vadd([1], [2], [0])
```

```
; SPIR-V
; Version: 1.0
; Bound: 45
; Schema: 0
```



```
OpCapability Addresses
OpCapability Linkage
OpCapability Kernel
...
OpStore %43 %39 Aligned 8
OpReturn
OpFunctionEnd
```

Level Zero wrappers

```
julia> using .oneL0
```

```
julia> drivers()
```

```
ZeDriver iterator for 1 drivers:
```

```
1. ZeDriver(00000000-0000-0000-174c-dd890103629a): version 1.3.25242
```

```
julia> drv = first(drivers());
```

```
julia> devices(drv)
```

```
ZeDevice iterator for 1 devices:
```

```
1. Intel(R) Arc(TM) A770 Graphics [0x56a0]
```

```
julia> dev = first(devices(drv));
```

```
julia> queue = ZeCommandQueue(ctx, dev);
```

```
julia> execute!(queue) do list
    append_barrier!(list)
end
```

Issues developing oneAPI.jl

- ✗ SPIR-V fragility
 - SPIRV-LLVM-Translator: incomplete LLVM support, invalid SPIR-V
 - Easy to trigger IGC aborts
- ✗ Math libraries (oneMKL, oneDNN, etc) are problematic
 - lack of C APIs
 - difficult to redistribute (cf. libcublas.so)
 - assumes SYCL environment (events, queues, etc)

Work in progress

- Performance: optimization for Intel hardware
- Integration with performance tools (VTune)
- Platform support: Linux & Intel GPUs only

Try it out!

<https://github.com/JuliaGPU/oneAPI.jl>