

Unified Runtime

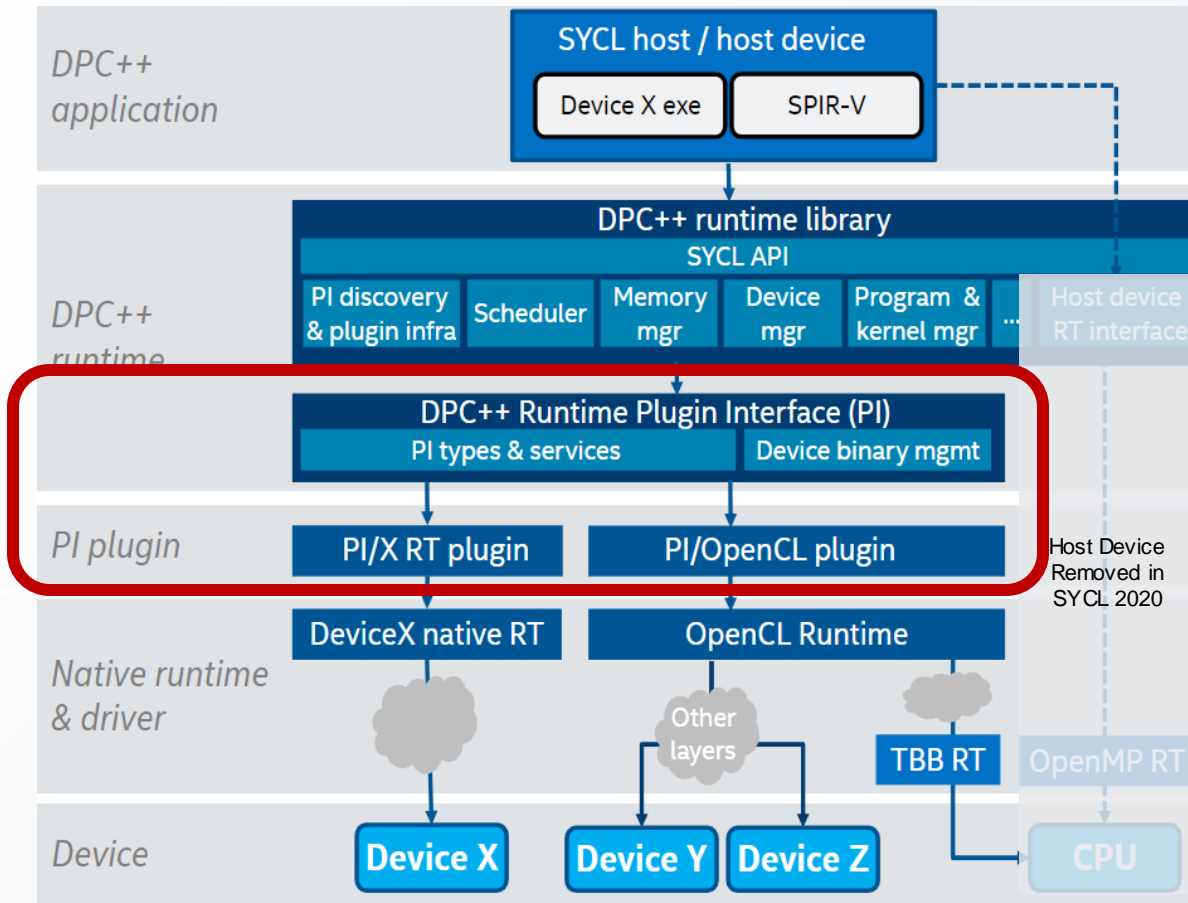


Topics

- Recap of Unified Runtime and goals
- Specification work
- Implementation work
- Unified Memory Architecture
- Future considerations

Specification

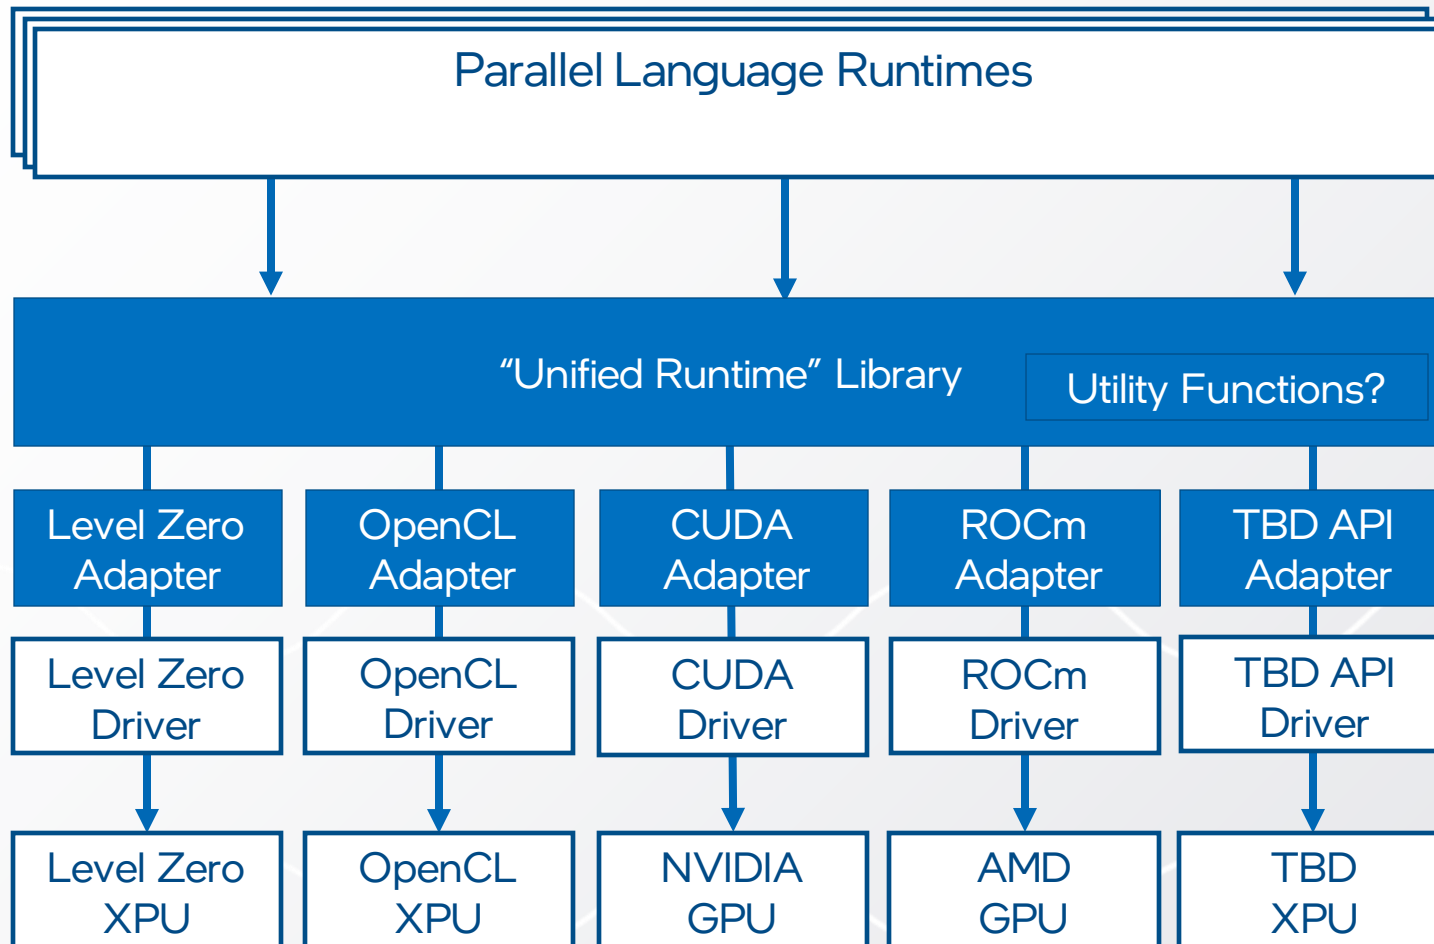
Recap: DPC++ Runtime Plugin Interface



Problem Statement:

- Plugin Interface is an implementation detail
 - No specification
- Plugin Interface is only usable by the DPC++ Runtime
 - Other language runtimes end up duplicating the same functionality
 - Interop between languages on separate implementations requires a lot of work

Discussion from Previous TABs: Unified Runtime



- Define a “Unified Runtime”, usable by any Parallel Language Runtime, with a well-defined interface.
- Reviewed draft specification:
 - <https://spec.oneapi.io/unified-runtime/latest/>
 - Previously 0.5, now 0.6.

Goals of Unified Runtime


- Provide portable glue layer between language runtimes and close-to-the-metal native device APIs
- Allow efficient implementations regardless of native API
- Provide common ease-of-use features needed by language runtimes
- Support interoperability and access to native backends
- Have reasonable implementation cost for runtime users and adapter developers
- Tool and debug support
- Next steps:
 - Short-term, upgrade DPC++ SYCL runtime and experiment with OpenMP
 - Over time, possible to earn wide engagement by open-source community


Specification v0.6

February 15, 2023 – March 15, 2023

Overview


78 Active pull requests

 62
Merged pull requests









 16
Open pull requests

Excluding merges, **10 authors** have pushed **87 commits** to main and **87 commits** to all branches. On main, **189 files** have changed and there have been **34,125 additions** and **17,206 deletions**.

- Specification work happens on GitHub directly
 - 192 PRs merged, 17 in review
 - 102 issues closed, 51 open
 - <https://github.com/oneapi-src/unified-runtime>
- Adapter and runtime porting in DPC++ codebase
 - <https://github.com/intel/llvm>
 - See PRs with “[UR]” tag
- A lot of specification changes are driven by keeping up with PI
 - DPC++, SYCL, and PI are always advancing, PI is not frozen, unified runtime matches each new PI change
- Other changes include mass rename to urFoo, large changes to compiler interface, changes to setting kernel arguments

Implementation

Experimental Implementation

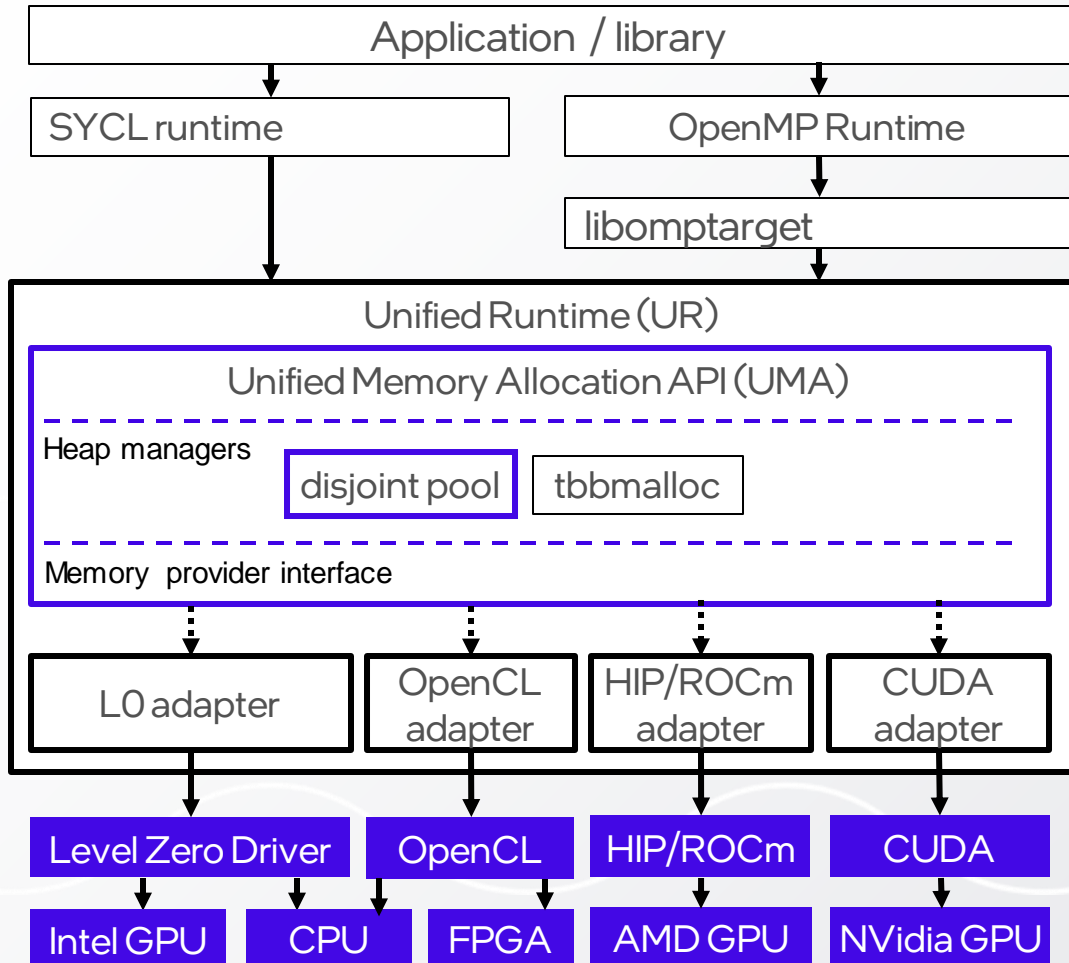

 cuda	[SYCL][CUDA][HIP]
 esimd_emulator	[SYCL] Move backe
 hip	[SYCL][CUDA][HIP]
 level_zero	[SYCL] Check that v
 opencl	[SYCL] Add device_
 unified_runtime	[SYCL][L0] Rework
 CMakeLists.txt	[SYCL][L0] POC for
 Id-version-script.txt	[SYCL] Export PI O



<https://github.com/intel/llvm/tree/sycl/sycl/plugins>

- Implement Unified Runtime as a PI plugin initially
- Existing PI plugins get rewritten as UR adapters
 - Per-adapter choice, but so far all are planning to base PI and UR adapters on common code for transition
 - Currently porting L0 & CUDA adapters
 - Next comes OpenCL & HIP
- Finally, once PI and UR have parity, refactor SYCL-RT to use UR directly

Unified Memory Architecture



- UR provides interfaces for memory allocations
- UR employs UMA for memory pooling
- UMA provides a unified interface for heterogeneous memory allocations
 - Decouple pool management from memory providers
- Adapters act as memory providers for coarse-grain allocations
- UMA is part of the upcoming UR release. Going to be a separate component in the next releases.

Future Considerations

License

- Current license on Unified Runtime repository is MIT
 - Covers common code and specification source, will cover adapters
- Project will involve taking code from DPC++ SYCL-RT, and wants long-term option of collaborating with upstream LLVM
 - All “Apache License v2.0 with LLVM Exceptions” license
 - So could change the license to match now, while project is small
- If this may negatively affect you please let us know
 - <https://github.com/oneapi-src/unified-runtime/issues/363>

Community Collaborations

- Previous TAB meetings asked us to look at OpenCL 3.0 and on-going work in LLVM libomptarget
 - Previous TAB in November covered OpenCL 3.0
- Short-term goal is to update SYCL runtime implementation, and investigate OpenMP
 - Learn more about how the API should look like through this
 - Improving the implementation and features of current compiler
- Bigger ambitions for the future
 - Produce widely used library and tooling for implementing language runtimes on a range of interfaces and hardware
 - Work with community on how best to achieve that

Summary

- Unified runtime specification and implementation are under current, active development
- Great time to give us any feedback!
 - Comments on recent specification changes: <https://spec.oneapi.io/unified-runtime/latest/>
 - Thoughts on building language runtimes on top of unified runtime
 - We're looking at SYCL and experimenting with OpenMP, but interested in language runtimes more broadly
 - Can file discussion points at <https://github.com/oneapi-src/unified-runtime/issues>

Questions