

SYCL-BLAS as a oneMKL backend

Ouadie EL FAROUKI

Staff Software Engineer

oneAPI Community Forum - Math Special Interest Group - July 12, 2023

Summary

- I. Introduction
- II. SYCL-BLAS
- III. SYCL-BLAS & oneMKL
- IV. Future plans
- V. Q&A

Introduction

BLAS

- **Basic Linear Algebra Subroutines** : Standardized in the late 70s/80s while trying to make it available & re-usable for everyone. Still in use today thanks to its **interface**.
- Main motives :
 - *Reusability through a common interface within higher level libraries & application.*
 - *Exploring hardware capabilities for efficiency & accuracy.*
- **BLAS** consists of 3 different levels of operations:
 - *level 1 : vector operations*
 - *level 2 : matrix - vector operations*
 - *level 3 : matrix - matrix operations*

Introduction

Who we are

Focus & missions :

- Expanding & maintaining SYCL BLAS library (*Operators, performance, portability etc..*).
- Building & maintaining CI & benchmarks of the SYCL BLAS library across different platforms.
- Integrating & maintaining SYCL-BLAS as a oneMKL backend.

Introduction

BLAS Libraries

- Many implementations have been developed as standalone libraries, each with a specific '*goal in mind*' :
 - **Open-source:** *OpenBlas, clBLAST, LAPACK*
 - **Proprietary:**
 - Intel MKL (*Intel CPUs & GPUs*)
 - cuBLAS (*NVIDIA GPUs*)
 - rocBLAS (*AMD GPUs*)
- The open-source implementations usually offer a functional portability along with the possibility of tuning a specific routine on specific platforms.
- The proprietary implementations usually offer fine-tuned performance for corresponding native hardware.

SYCL-BLAS

Overview

- It's a SYCL & C++ based BLAS implementation started in 2015.
- It aims to be the reference BLAS implementation for tuneable performance and **portability**, as a community open-source project.
- Many papers have been published about it :
 - Aliaga, José I., Ruymán Reyes, and Mehdi Goli. *"SYCL-BLAS: leveraging expression trees for linear algebra."* Proceedings of the 5th International Workshop on OpenCL. 2017.
 - Aliaga, José I., Ruymán Reyes, and Mehdi Goli. *"SYCL-BLAS: combining expression trees and kernel fusion on heterogeneous systems."* Parallel Computing is Everywhere 32 (2018): 349.
 - Sabino, Thales, and Mehdi Goli. *"Toward Performance Portability of Highly Parametrizable TRSM Algorithm Using SYCL."* International Workshop on OpenCL. 2021.

About

An implementation of BLAS using the SYCL open standard for acceleration on OpenCL devices

📖 Readme

📄 Apache-2.0 license

🤝 Code of conduct

📈 Activity

★ 175 stars

👁 24 watching

🍴 45 forks

github.com/codeplaysoftware/sycl-blas/

SYCL-BLAS

Overview

- SYCL-BLAS Follows modern C++ specifications :
 - Can be used as Header-only library (*no linking needed, simple `#include<sycl_blas.h>`*)
 - Uses template meta-programming for maximum flexibility and performance (*Expression trees similar to Eigen's approach*).
- Developed to be the reference SYCL-Based implementation :
 - It can run on any SYCL compatible device (*through supporting SYCL implementation*).
 - It can be compiled with different SYCL compilers :
 - *DPC++ / icpx*
 - *ComputeCPP*
 - *HIPSYcl (OpenSYCL)*

SYCL-BLAS

Performance & Portability

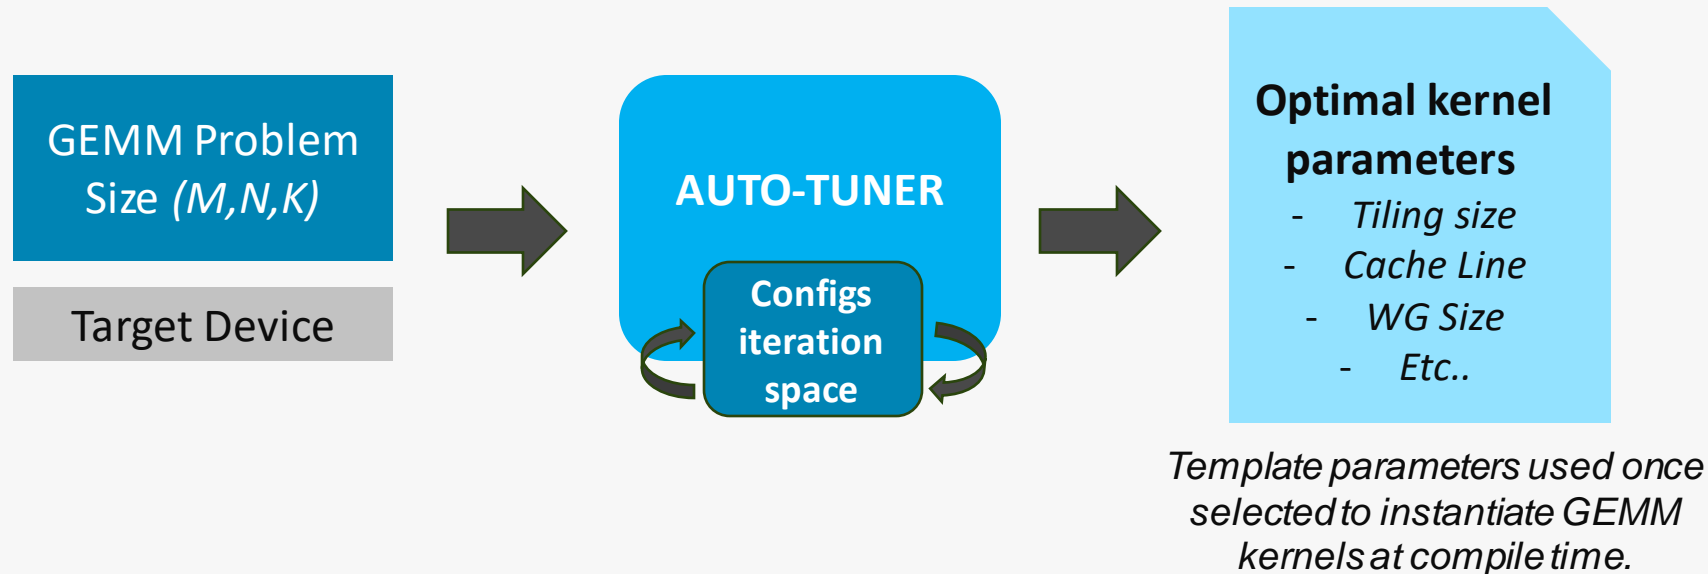
- SYCL BLAS Operators/Kernels support & can be tuned for specific targets thanks to templated parametrizations :
 - *NVIDIA GPUs*
 - *AMD GPUs*
 - *INTEL GPUs*
 - *Default CPU*

```
~ >> cmake -GNinja .. \  
-DSYCL_COMPILER={dpcpp | computecpp | hipsycl .. } \  
-DTUNING_TARGET={DEFAULT_CPU | NVIDIA_GPU | ..} \  
-DDPCPP_SYCL_TARGET={spir64 | nvptx64-nvidia-cuda | ..}  
-DDPCPP_SYCL_ARCH={sm_XY | gfxXYZ .. }
```


SYCL-BLAS

Auto-tuning : GEMM

- Automatically calculate the optimal parameters for GEMM on the given platform, maximizing performance*.

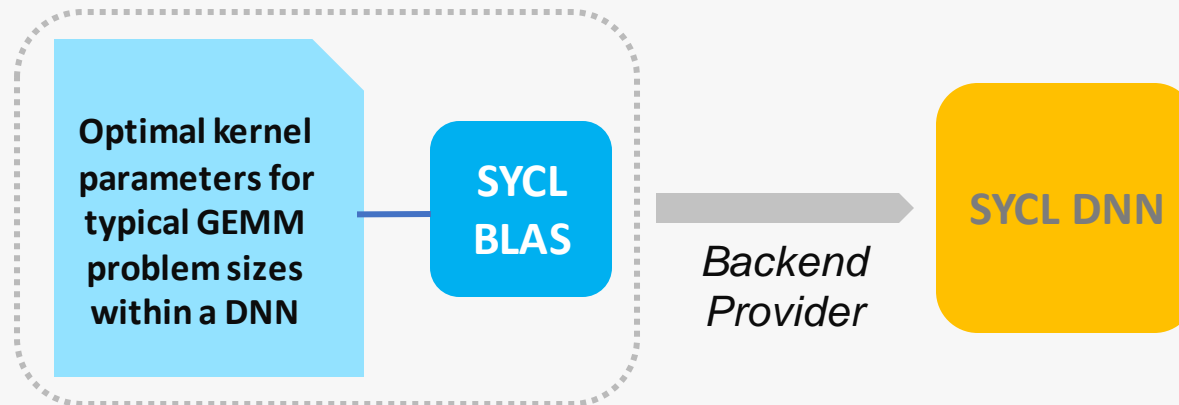


*: github.com/codeplaysoftware/sycl-blas/tools/auto_tuner

SYCL-BLAS

Some ongoing use-cases

- SYCL-DNN (A Deep Neural Networks SYCL Library) exposes an optional **SYCL-BLAS backend** to be used for some operations (e.g., Memory handling, Matrix Multiplication, Reduction etc..).

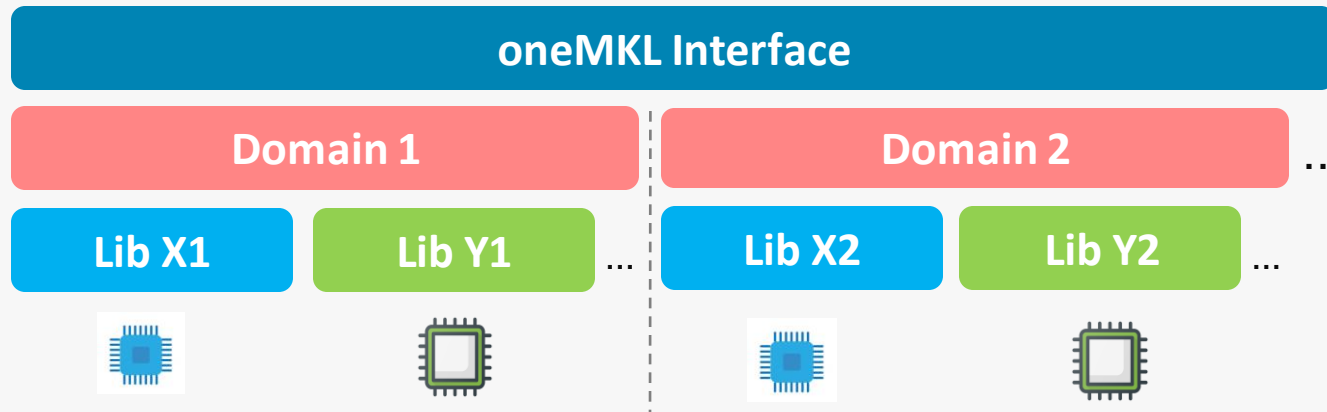


- A SYCL Based *JPEG-Compression* app has been developed using SYCL-BLAS Gemm as a **tunable building block** for performance portability.

OneMKL

Overview

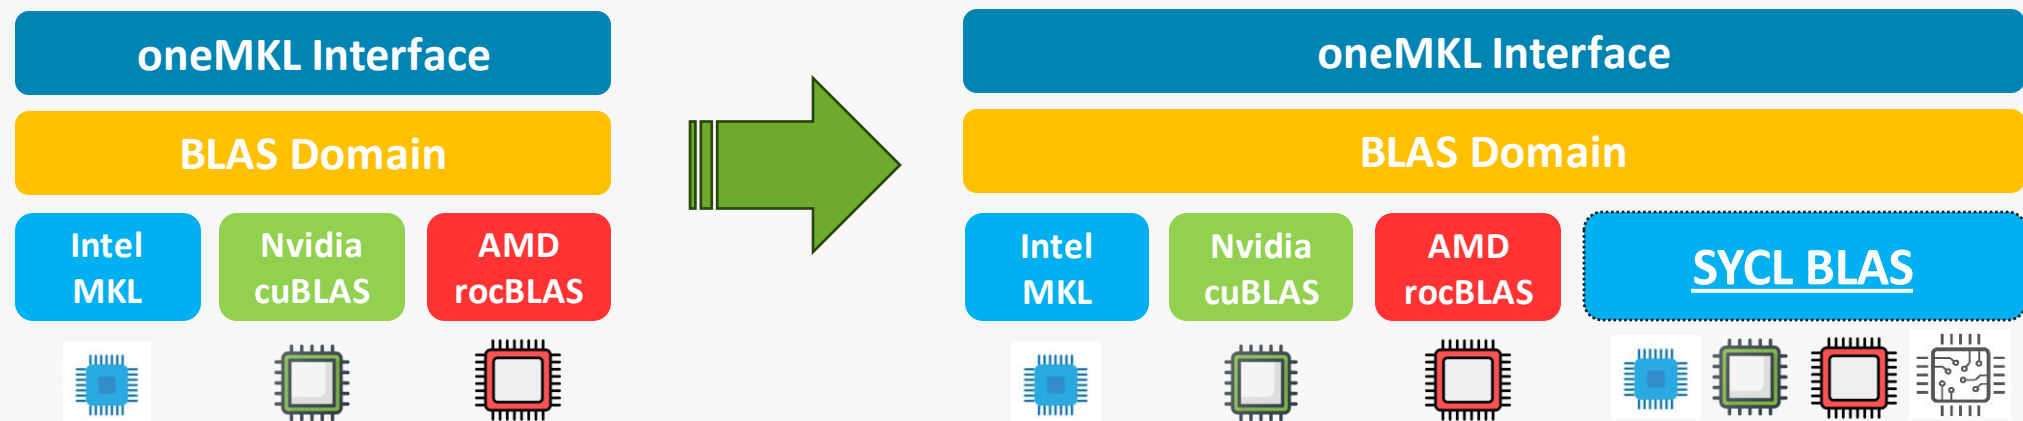
- oneAPI Math Kernel Library :
 - Open-source implementation of the DPC++ oneMKL Spec & part of the oneAPI 10 core specs.
 - It's a library gathering optimized common mathematical routines grouped into '*domains*' : **BLAS**, **LAPACK**, **RNG**, **DFT** etc..
 - It supports multiple devices through relevant libraries (**backends**) underneath.



SYCL-BLAS & oneMKL

SYCL-BLAS Backend

- For the BLAS domain in oneMKL, there are backends relying on proprietary third-party libraries (*MKL*, *cuBLAS*, *rocBLAS*), each supporting a set of native devices.
- Introducing SYCL-BLAS as a BLAS backend to oneMKL aims to support portability in *oneMKL* interface within the backend itself while remaining open-source :



SYCL-BLAS & oneMKL

SYCL-BLAS Backend

- To use SYCL-BLAS as backend, users don't need to install it by themselves. It's taken care of by the built-in configurations.
- Use local version of *SYCL-BLAS* by specifying target directory.

```
~ >> cmake -GNinja .. \  
        -DENABLE_SYCLBLAS_BACKEND=ON \  
        -DENABLE_MKLCPU_BACKEND=OFF \  
        -DENABLE_MKLGPU_BACKEND=OFF \  
        -DTARGET_DOMAINS=blas
```

```
~ >> cmake -GNinja .. \  
        -DENABLE_SYCLBLAS_BACKEND=ON \  
        -DENABLE_MKLCPU_BACKEND=OFF \  
        -DENABLE_MKLGPU_BACKEND=OFF \  
        -DTARGET_DOMAINS=blas \  
        -Dsycl_blas_DIR=[INSTALLED DIR  
OF SYCLBLAS headers]
```

SYCL-BLAS & oneMKL

Usage Model

- **Run-time dispatching** : The application is linked with the oneMKL library, and the backend is loaded at run-time based on device vendor.

```
#include "oneapi/mkl.hpp"

gpu_dev = sycl::device(sycl::gpu_selector_v);

sycl::queue gpu_queue(gpu_dev);

oneapi::mkl::blas::column_major::gemm(gpu_queue, transA, transB, m, ...);
```

SYCL-BLAS & oneMKL

Usage Model

- **Compile-time dispatching** : The application used a templated backend selector API where the template parameters specify the required backend and third-party libraries. The application is linked with the required oneMKL backend wrapper libraries.

```
#include "oneapi/mkl.hpp"

gpu_dev = sycl::device(sycl::gpu_selector_v);
sycl::queue gpu_queue(gpu_dev);

oneapi::mkl::backend_selector<oneapi::mkl::backend::syclblas>
gpu_selector(gpu_queue);

oneapi::mkl::blas::column_major::gemm(gpu_selector, transA, transB, ...);
```

SYCL-BLAS & oneMKL

Example

Device Selection	{	<code>auto gpu_dev = sycl::device(sycl::gpu_selector_v);</code>
		<code>sycl::queue gpu_queue(gpu_dev);</code>
Backend Selection	{	<code>oneapi::mkl::backend_selector<oneapi::mkl::backend::syclblas> gpu_selector(gpu_queue);</code>
		<code>[...]</code>
Preparing buffers & launching kernel	{	<code>{</code>
		<code>auto buffer_a = sycl::buffer{A.data(), sycl::range{matrix_size}};</code>
		<code>auto buffer_b = sycl::buffer{B.data(), sycl::range{matrix_size}};</code>
		<code>auto buffer_c = sycl::buffer{C.data(), sycl::range{matrix_size}};</code>
		<code>oneapi::mkl::blas::column_major::gemm(gpu_selector, transA, transB, m, n, k,</code>
		<code>alpha, buffer_a, lda, buffer_b, ldb,</code>
		<code>beta, buffer_c, ldc);</code>
		<code>}</code>

Future plans

SYCL-BLAS is in active development and there are more upcoming features.

- USM full support.
- Row-major support.
- Increase operator implementation (*currently ~60%*).
- Complex type support.

SYCL-BLAS is the portable backend of oneMKL.

Update


Name changing of the library to align with *oneAPI*'s strategy & emphasize its main feature : **portability**.

SYCL-BLAS

Update

Name changing of the library to align with *oneAPI*'s strategy & emphasize its main feature : **portability**.

portBLAS



Q&A



@codeplaysoft



info@codeplay.com



codeplay.com

Notices & Disclaimers

Intel technologies may require enabled hardware, software or service activation.

No product or component can be absolutely secure.

Your costs and results may vary.

© Codeplay Software Ltd.. Codeplay, Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.