



oneAPI Level Zero Technical Advisory Board

Welcome and Thanks

- A unique opportunity to steer the parallel programming ecosystem
- A problem worth solving
 - Multi-architecture, avoiding lock-in to 1 specific hardware architecture
 - Direct and library-based programming
 - Extending existing models
 - Performant
- Your leadership, input, and feedback is critical

Rules of the Road

- DO NOT share any confidential information or trade secrets with the group
- DO keep the discussion at a High Level
 - Focus on the specific Agenda topics
 - We are asking for feedback on features for the oneAPI specification (e.g. requirements for functionality and performance)
 - We are NOT asking for feedback on any implementation details
- Please submit any feedback in writing on Github in accordance with the [Contribution Guidelines](https://spec.oneapi.com/contribution-guidelines) at spec.oneapi.com. This will allow Intel to further upstream your feedback to other standards bodies, including The Khronos Group SYCL* specification.

Agenda

- oneAPI Welcome & Introduction – Paul Petersen, Intel
- Level Zero:
 - Specification & How to Participate -- Zack Waters, Intel
- Discussion Topic:
 - Separation of Sysman from core Level Zero APIs -- Ben Ashbaugh, Intel
- Wrap up, Question & Answer – All

oneAPI

A unified programming model
to simplify development
across diverse architectures

Common developer experience across architectures

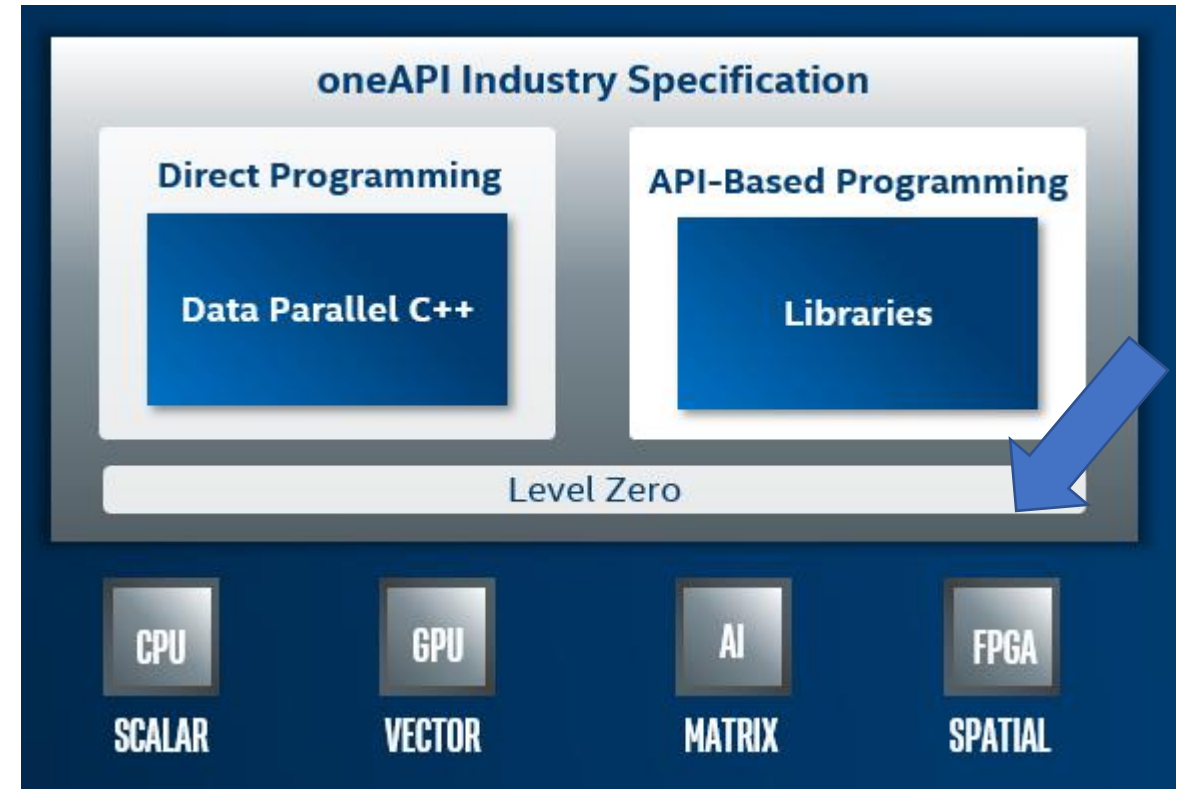
Unified and simplified language and libraries for
expressing parallelism

Uncompromised native high-level language
performance

Interoperates with existing languages and libraries

Support for CPU, GPU, AI and FPGA

Based on industry standards and open specifications

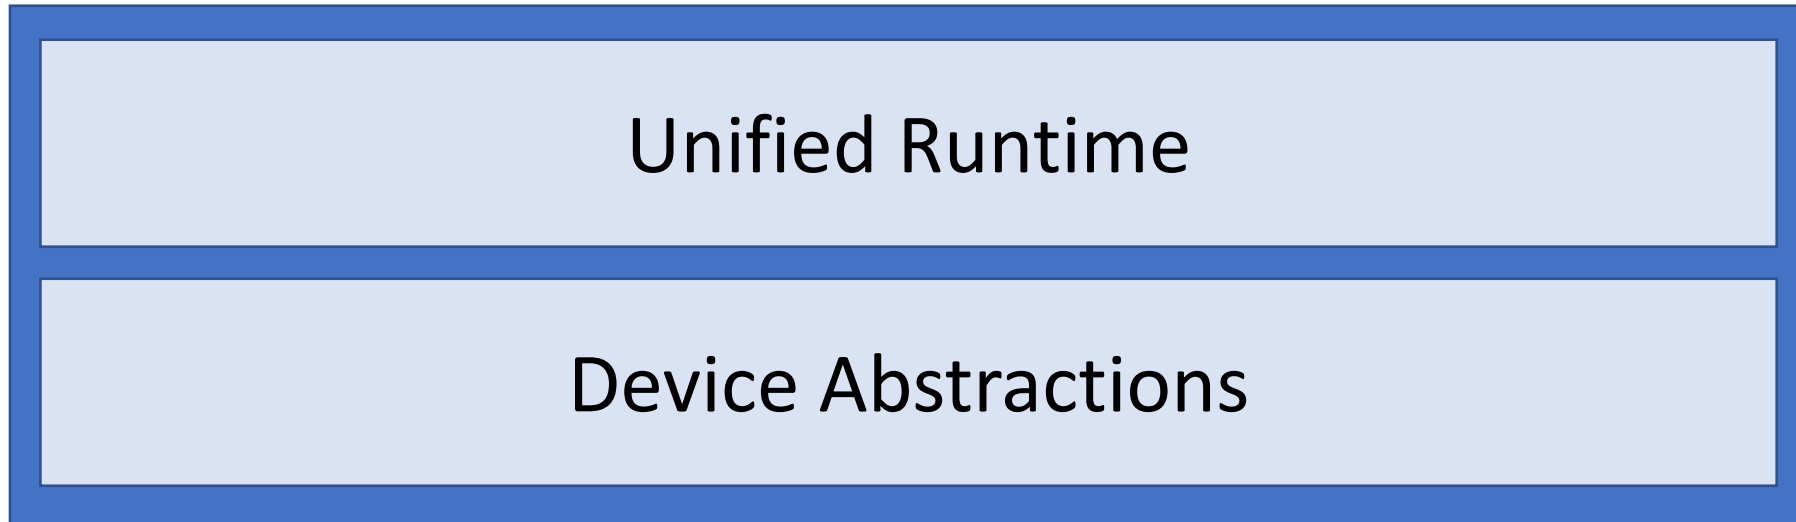


Does oneAPI need a Unified Runtime?

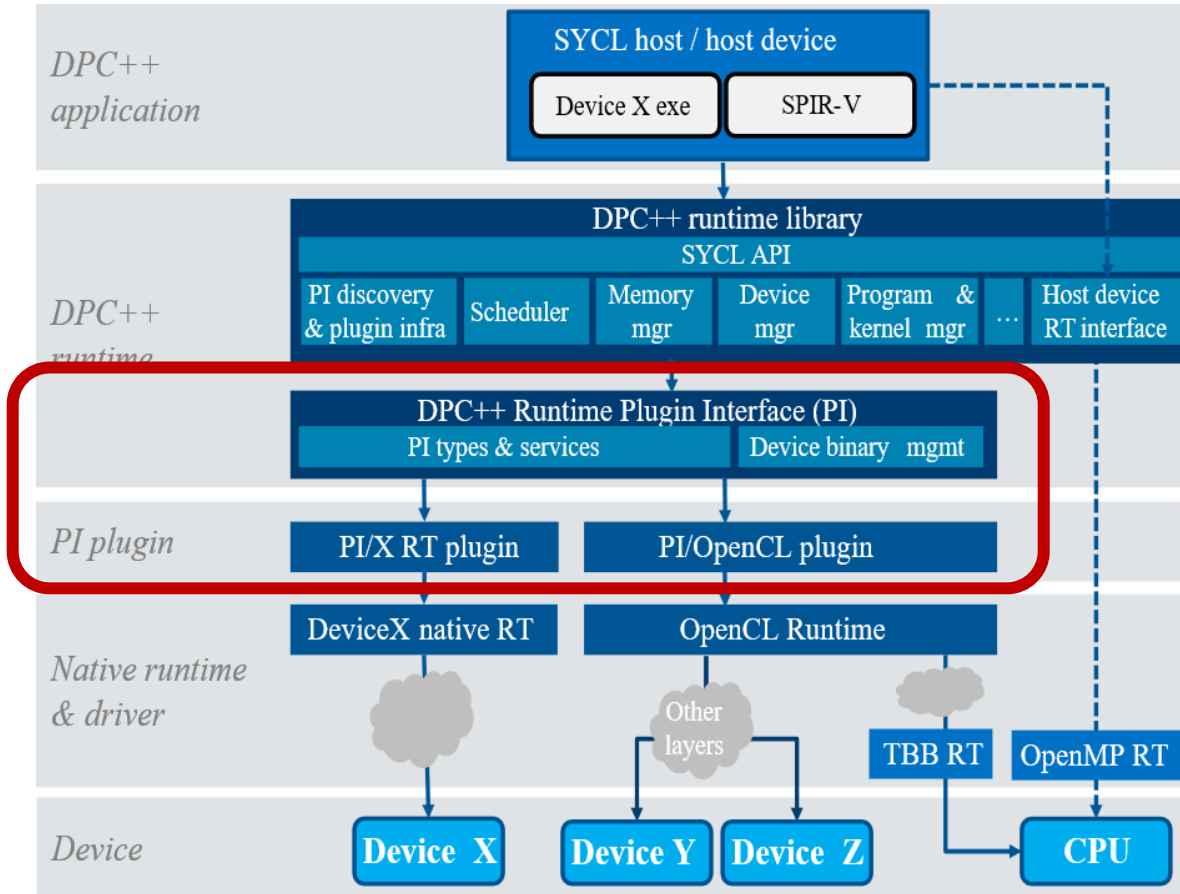
- Why is Level Zero as a defined not already sufficient?
 - Level Zero provides a low-level abstraction to a device.
 - Can we provide richer constructs for cross-platform and cross-API enabling?
- What are the [Design principles](#)?
 - Is it Deterministic – How would you support tuning heuristics?
 - Is it Stateless – How would you support callbacks to be registered?
- What features are missing to support your favorite language?
- Do we need special support for the CPU?
 - CPU device driver or resource management?
 - Leverage common components -- [hwloc](#) / [memkind](#) / [numa](#)?

Evolution of Level Zero

Level Zero



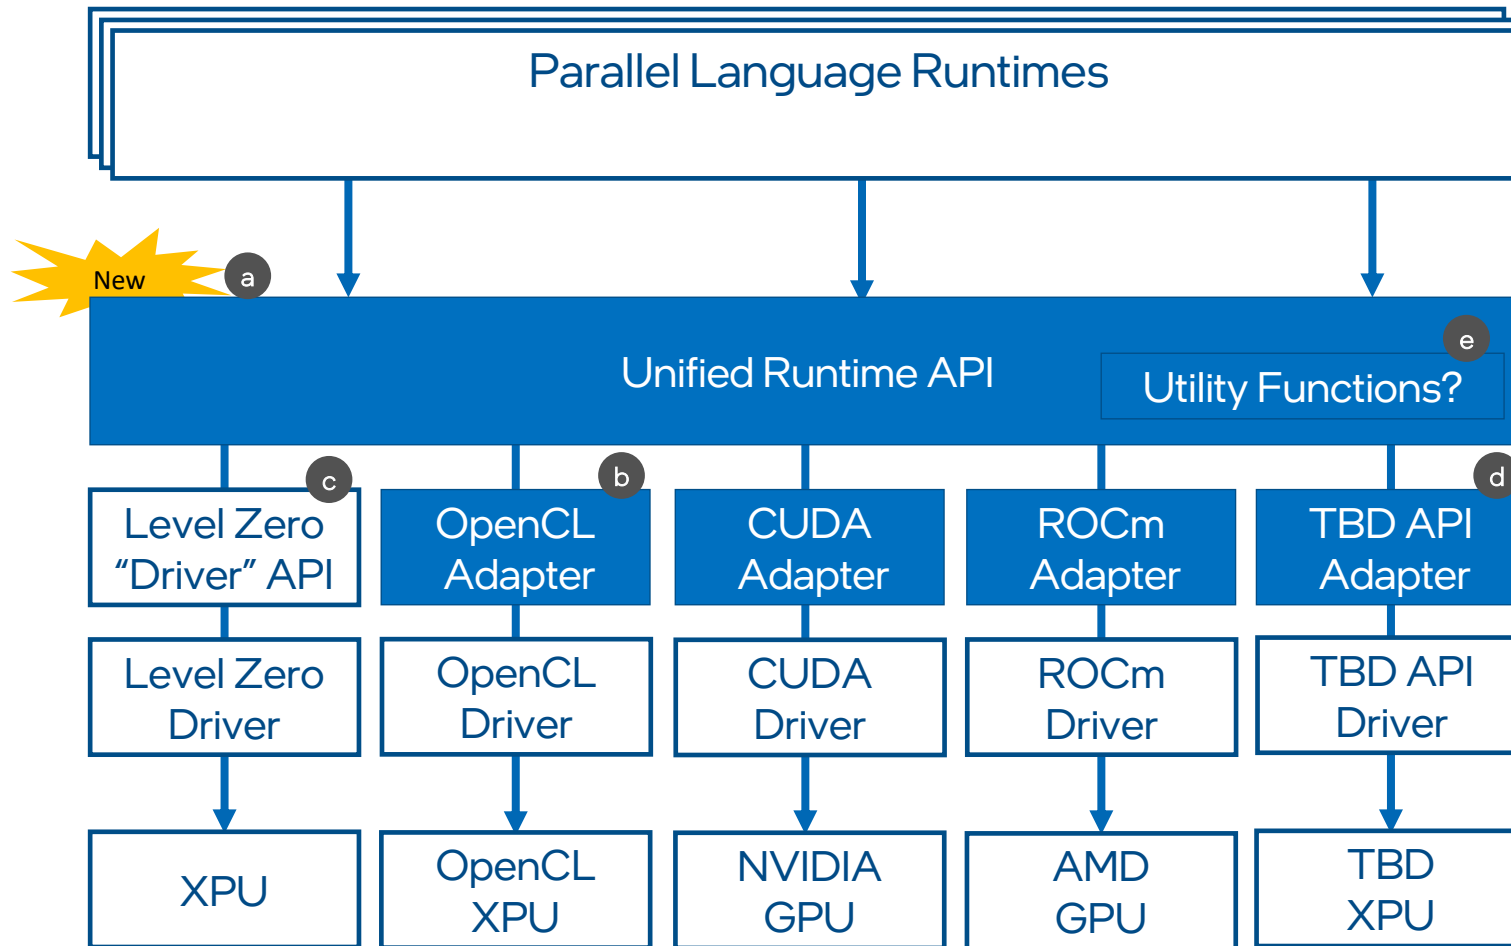
DPC++ Runtime Plugin Interface



Problem Statement:

- Plugin Interface is an implementation detail
 - No formal specification
- Plugin Interface is only usable by the DPC++ Runtime
 - Other language runtimes must duplicate functionality
- Plugin Interface has deep OpenCL heritage

Proposal: Unified Runtime API




- a. Refactor and formalize Plugin Interface into Unified Runtime APIs for use by multiple language runtimes
- b. Refactor existing plugins into Unified Runtime API Adapters
- c. Refactor, generalize, and modernize adapter interface
- d. Enable new backends by implementing new Unified Runtime API Adapter
- e. Eventually: Move common utility functionality to Unified Runtime API?

Request

“Our” initiative – help to frame it for your needs

Be brutally honest – good and bad

Invite your friends – and have fun



Thank You!

<http://oneapi.com>